

BY FENIL GAJJAR

KUBERNETES DAILY TASKS

KUBERNETES TAINTS TOLERATIONS & NODE SELECTORS

- COMPREHENSIVE GUIDE
- THEORY + PRACTICAL TASKS
- REAL TIME SCENARIO TASKS
- END TO END DOC

100%
SOLID

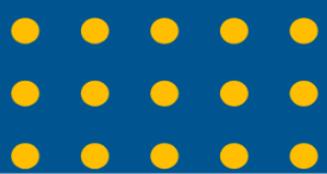


CONTACT US

fenilgajjar.devops@gmail.com

linkedin.com/in/fenil-gajjar

[gitub.com/Fenil-Gajjar](https://github.com/Fenil-Gajjar)





Mastering Kubernetes: Taints, Tolerations, Labels & Selectors | From Core Concepts to Real-World Deployments



Welcome!

Hey there!

Thank you so much for stopping by and exploring this detailed doc on **Kubernetes Taints, Toleration, Node Labels & Selectors.** 

Whether you're a DevOps engineer, SRE, platform engineer, or a curious learner diving into Kubernetes infrastructure, you're in the right place! 

About This Doc

This isn't just theory — this guide is a **deep dive from fundamentals to production-grade use cases**. We'll walk through:

-  **Understanding** what taints, tolerations, labels, and selectors really do under the hood.
-  **How they work together** to influence pod scheduling, isolation, and workload control.
-  **Real-world implementation tasks** using `kops`-based Kubernetes clusters.
-  **Troubleshooting tips** and  **best practices** you can apply in enterprise-grade clusters.

If you're managing a multi-tenant cluster, deploying ML workloads, isolating nodes for platform services, or just trying to cleanly control where your pods land — this doc is for you.



Part of a Bigger Series

This document is part of my growing DevOps content series where I share hands-on knowledge, practical scenarios, and production patterns across:

- **AWS**
- **Ansible**
- **Kubernetes**
- **Infrastructure as Code**
- And more cloud-native tooling and practices

You can follow my content, series updates, code samples, and more on:

LinkedIn: linkedin.com/in/fenil-gajjar

GitHub: github.com/Fenil-Gajjar

 Stay Tuned!

Thank you again for reading — and **extra gratitude** if you've checked out my previous docs as well. Your support means a lot! 🙏

 **More deep dives are coming soon**, so stay connected and get ready for more content just like this — clean, practical, and made for real DevOps engineers.

Let's keep learning and building together.

Happy Clustering!  



What Are Taints and Tolerations in Kubernetes?

In Kubernetes, **taints** and **tolerations** work together to **control which pods can be scheduled on which nodes**. This mechanism provides **fine-grained control** over **pod placement** by allowing certain nodes to repel pods unless those pods explicitly "tolerate" the condition.

- ◆ **Taints – Think of them as Node-Level Restrictions**

A **taint** is applied to a node. It marks the node as **unsuitable for general scheduling**. Once a node is tainted, **no pods** will be scheduled onto it **unless they have a matching toleration**.

A taint consists of three parts:

`<key>=<value>:<effect>`

Example:

`key=env, value=prod, effect=NoSchedule`

This means:

Don't schedule any pod on this node unless it tolerates a taint with key **env**, value **prod**, and effect **NoSchedule**.

- ◆ **Tolerations – Think of them as Pod-Level Permissions**

A **toleration** is applied to a **pod**. It tells Kubernetes:

“I am okay with being scheduled on a node that has a matching taint.”

Tolerations don’t force scheduling onto tainted nodes. They just allow it.

YAML example:

tolerations:

```
- key: "env"  
  operator: "Equal"  
  value: "prod"  
  effect: "NoSchedule"
```

- ◆ **Effects of Taints**

There are **three types of effects** that define how Kubernetes should treat a taint:

Effect	Behavior
---------------	-----------------

NoSchedule	Do not schedule pods on the node unless they tolerate the taint
-------------------	---

PreferNoSch Try to avoid scheduling pods here, but not strictly
edule enforced

NoExecute Existing pods that don't tolerate this taint will be
evicted

◆ Use Cases (When and Why)

- **Dedicated nodes:** Isolate workloads (e.g., run only GPU workloads or sensitive apps).
- **Environment isolation:** Separate **dev**, **test**, and **prod** workloads.
- **Cluster maintenance:** Temporarily cordon nodes for draining or updates.
- **Spot instances:** Taint ephemeral/spot instances to prevent critical workloads from landing there.

Why Do We Need Taints and Tolerations in Kubernetes?

Kubernetes by default tries to **utilize all nodes evenly** and will schedule pods wherever there are available resources. But in real-world clusters, **not all nodes are created equal**—and **not all workloads should run everywhere**.

That's where **taints and tolerations** come in. They give platform engineers the ability to **control where pods can and cannot run**. This is **not a resource constraint**, but rather a **policy constraint**.

Without Taints and Tolerations:

Imagine a cluster where:

- Some nodes are **GPU-enabled** for ML workloads.
- Some nodes are **spot instances**, cheaper but **less reliable**.
- Some nodes are **reserved for production traffic only**.

If there's no control:

- Regular, low-priority workloads could land on **expensive GPU nodes**.
- Critical production pods could be scheduled on **spot nodes** and lost.
- Dev/test workloads might run on **production-isolated** nodes.

This leads to:

- **Wasted resources**
- **Performance issues**
- **Security/risk exposure**
- **Unpredictable behavior**

With Taints and Tolerations:

Taints let you **mark nodes** for specific purposes.

Tolerations let **only approved pods** be scheduled on those nodes.

This results in:

- **Workload isolation** (e.g., dev vs. prod)
- **Node specialization** (e.g., GPU, memory-optimized, ARM nodes)
- **Improved reliability** (e.g., avoid unreliable nodes)
- **Policy enforcement** without tight coupling to pod specs

How Do Taints and Tolerations Work in Kubernetes?

At a high level, **taints and tolerations work together** to allow **node-level filtering**. They don't assign pods to specific nodes — instead, they define **where pods are not allowed** to run, unless they **tolerate** the restriction.

Let's break it down in detail:

◆ 1. Taints Are Applied on Nodes

A **taint** is a property you attach to a **node** to indicate that it should repel all pods — **except those that can tolerate it**.

The taint uses this format:

```
<key>=<value>:<effect>
```

Example:

```
kubectl taint nodes node-1 environment=prod:NoSchedule
```

This command applies a taint with:

- **key: environment**
- **value: prod**

- **effect: NoSchedule**

Effect types:

Effect	Description
NoSchedule	Pods that do not tolerate the taint will not be scheduled on the node.
PreferNoSchedule	Kubernetes will try to avoid scheduling non-tolerating pods on the node, but won't strictly block them .
NoExecute	Not only blocks scheduling — it evicts running pods that don't tolerate the taint.

◆ 2. Tolerations Are Applied on Pods

A **toleration** is specified in the pod's manifest. It signals that the pod **understands and accepts** the taint on a node and is willing to run there.

Example YAML:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

spec:

 containers:

 - name: nginx

 image: nginx

 tolerations:

 - key: "environment"

 operator: "Equal"

 value: "prod"

 effect: "NoSchedule"

This pod will **tolerate the taint** we applied to **node-1**.

Important note:

Tolerations don't force Kubernetes to place a pod on a tainted node — they **only allow it**.

🔍 How Kubernetes Uses This Mechanism Internally

When Kubernetes schedules a pod:

1. It looks at **all available nodes**.
2. It filters out any node with a taint that the pod **does not tolerate**.

-
3. From the remaining nodes, it uses other scheduling policies (like resource availability, affinity, etc.) to pick a node.

So:

- **No matching toleration → pod excluded from tainted node**
- **Matching toleration → node considered as a valid candidate**



Example: Combined Flow

Let's say we have:

Node A:

```
kubectl taint nodes node-a app=ml:NoSchedule
```

Pod:

```
tolerations:
```

```
- key: "app"  
  operator: "Equal"  
  value: "ml"  
  effect: "NoSchedule"
```

👉 Result: Pod **can** be scheduled onto **node-a**.

⚠️ Using **NoExecute** with **TolerationSeconds**

NoExecute is a powerful taint effect. It doesn't just block new pods — it will **evict** **running ones** if they don't tolerate it.

You can also specify how **long** a pod is allowed to stay after the taint appears.

tolerations:

```
- key: "maintenance"  
  operator: "Equal"  
  value: "true"  
  effect: "NoExecute"  
  
  tolerationSeconds: 300
```

➡️ This pod will stay on the tainted node for 300 seconds, then be evicted.

🔧 How to Add/Remove Taints

Add a taint:

```
kubectl taint nodes node-1 key=value:effect
```

Remove a taint:

```
kubectl taint nodes node-1 key:effect-
```

Example:

```
kubectl taint nodes node-1 environment=prod:NoSchedule
```

```
kubectl taint nodes node-1 environment:NoSchedule- #
```

Removes it

 **Operator Field in Tolerations**

The **operator** field controls how the key/value pair is interpreted.

Operator	Description
Equal	Key and value must both match
Exists	Only the key needs to exist (value ignored)

Example with Exists:

tolerations:

```
- key: "dedicated"
```

```
operator: "Exists"
```

```
effect: "NoSchedule"
```

This pod tolerates **any taint** with key **dedicated** — regardless of value.

How Does Scheduling Work in Kubernetes (in the Context of Taints and Toleration)?

To answer it simply:

Tolerations alone do *not* guarantee that a pod will be scheduled onto a tainted node.

They only **allow** scheduling onto that node — they **don't force** or **prioritize** it.

Let's break it down in stages:

◆ 1. The Scheduler's Job

Kubernetes has a **default scheduler** that is responsible for:

- Filtering nodes where a pod *can run*
- Scoring the remaining nodes
- Picking the most appropriate one based on scores

The scheduler evaluates:

- Taints & tolerations
- Node affinity / anti-affinity
- Resource requests (CPU/memory)
- NodeSelector or node pools
- Topology spread constraints
- Custom scheduler logic, if configured

◆ 2. Filtering Stage: Taints Come First

During the **filtering** phase, taints play a crucial role.

- If a node has a taint that a pod **does not tolerate** → ✅ Node is *excluded* from the candidate list.
- If a node has a taint that a pod **does tolerate** → ✅ Node is *included* in the candidate list.
- Nodes **without taints** → ✅ Always included.

👉 Taints are a **gatekeeper**. Tolerations are like **keys** that open those gates.

◆ 3. Scoring Stage: Toleration Are NOT a Factor

Once Kubernetes has a filtered list of nodes where the pod *can* be scheduled:

- It assigns scores to those nodes based on resource availability, topology, affinity rules, etc.
- **Taints/tolerations have no impact here.**
- The pod could still be scheduled on a *non-tainted* node if it scores higher.

✗ So What Doesn't Happen?

If a pod has a toleration for a taint:

- Kubernetes **won't prefer** tainted nodes.
- Kubernetes **won't guarantee** that the pod will run there.
- It might never run on the tainted node **unless no better nodes are available.**

This often surprises people.

✓ How to Ensure a Pod Runs on a Tainted Node?

If you want a pod to be *specifically scheduled* onto a tainted node, you **combine**:

1. **Taint on the node**
2. **Toleration on the pod**
3. **NodeSelector or node affinity** on the pod to *target* the desired node

Example: Guaranteed Scheduling onto a Tainted Node

Taint the node:

```
kubectl taint nodes gpu-node gpu=true:NoSchedule
```

Pod manifest:

```
spec:  
  nodeSelector:  
    gpu: "true"  
  
  tolerations:  
  - key: "gpu"  
    operator: "Equal"  
    value: "true"  
  
  effect: "NoSchedule"
```

Now the pod can *only* run on nodes with label `gpu=true`, and it tolerates the taint → **it will be scheduled there.**

Why Workloads Don't Get Scheduled on Control Plane

Nodes by Default?

Yes, **Kubernetes control plane nodes come with a default taint**, which **prevents regular workloads from being scheduled** on them.

This is a deliberate architectural choice to **protect the stability and performance of the cluster**.

Deep Dive: The Justification

The **control plane** is the **brain of the Kubernetes cluster**. It runs critical components like:

- `kube-apiserver`
- `etcd` (Kubernetes' database)
- `kube-scheduler`
- `kube-controller-manager`

-
- Cloud controller and DNS (in many setups)

These components must always be **highly available, lightweight, and protected**.

If user workloads were to be scheduled on the same nodes:

- Resource starvation (CPU/memory) could bring down `etcd` or `kube-apiserver`
- Cluster-wide instability could follow
- Even a misbehaving pod could block access to the whole cluster

That's why Kubernetes **by design** applies a **default taint** to control plane nodes.



The Default Taint on Control Plane Nodes

On most Kubernetes installations (including `kubeadm`, `kops`, EKS, etc.), control plane nodes have the following default taint:

`node-role.kubernetes.io/control-plane=:NoSchedule`

Or, in older setups:

`node-role.kubernetes.io/master=:NoSchedule`

This taint:

- Has no key or value
- Has **NoSchedule** effect
- Repels any pod **unless it has a matching toleration**

How You Can Verify This

Run:

```
kubectl describe node <control-plane-node-name>
```

You'll see a taint like:

Taints:

```
node-role.kubernetes.io/control-plane:NoSchedule
```

Or with `kubectl get nodes -o json`:

```
"taints": [  
  {  
    "effect": "NoSchedule",  
    "key": "node-role.kubernetes.io/control-plane"
```

```
}
```

```
]
```



Can You Schedule Pods on Control Plane Nodes?

Yes — but you must **explicitly override** the taint.

To do that, the pod must include a **toleration** for the taint:

tolerations:

```
- key: "node-role.kubernetes.io/control-plane"  
  effect: "NoSchedule"  
  
operator: "Exists"
```

But even with this toleration, **you also need to steer the pod** to that node (since Kubernetes will not *prefer* tainted nodes). So you typically also use:

- **nodeSelector**
- or **nodeName**

Example:

```
spec:  
  tolerations:  
    - key: "node-role.kubernetes.io/control-plane"  
      effect: "NoSchedule"  
      operator: "Exists"  
  
  nodeSelector:  
    node-role.kubernetes.io/control-plane: ""
```

⚠ Should You Do This in Production?

Strongly discouraged in production environments unless:

- You're on a **very small dev/test cluster**
- Or running **high-availability control planes** with sufficient isolation

In production:

- Keep control plane clean
- Keep workloads off it
- Use taints + separate node pools to control workload placement

🎯 Taint Effects in Kubernetes

Before diving into each, here's a quick reference of the three effects:

Effect Type	What it Does
NoSchedule	Prevents scheduling of non-tolerating pods on a node
PreferNoSch edule	Tries to avoid scheduling, but allows it if necessary
NoExecute	Prevents scheduling and evicts already running non-tolerating pods

Now, let's go deep into:

🔍 NoSchedule — Deep Dive

📌 What is NoSchedule?

The **NoSchedule** effect is used to **strictly prevent pods from being scheduled on a node**, unless those pods **explicitly tolerate** the taint.

- If a pod does *not* tolerate the taint → it will *never* be scheduled on that node.

When to Use **NoSchedule**

Use **NoSchedule** when you want to **enforce a hard boundary** between workloads and nodes. This is the most deterministic and strict form of scheduling exclusion.

Common real-world examples:

- Keep dev workloads off prod nodes
- Ensure only GPU-intensive apps run on GPU nodes
- Prevent pods from landing on spot or preemptible instances unless intended
- Reserve nodes for specific tenants or environments

Syntax and Example

Taint a node:

```
kubectl taint nodes node-a dedicated=prod:NoSchedule
```

This means:

 "Only pods that tolerate **dedicated=prod** with effect **NoSchedule** can be scheduled here."

Now the pod needs this toleration to be allowed:

tolerations:

- key: "dedicated"
operator: "Equal"
value: "prod"
effect: "NoSchedule"

→ This pod **now tolerates** the taint and *may* be scheduled on that node.

✗ What Happens if Pod Lacks the Toleration?

- The Kubernetes scheduler sees the node
- Filters it **out** from consideration
- Pod won't be scheduled there — **ever**

No retries, no random behavior — this is a **hard rule**.

Real-World Scenario Example

Problem:

You want to reserve certain nodes (`node-group-prod`) **strictly for production traffic only.**

Solution:

Taint the nodes:

```
kubectl taint nodes node-group-prod  
dedicated=prod:NoSchedule
```

Add toleration to prod-only workloads:

tolerations:

```
- key: "dedicated"  
  operator: "Equal"  
  value: "prod"  
  
  effect: "NoSchedule"
```

(Optional) Add a nodeSelector:

nodeSelector:

1. `dedicated: prod`

PreferNoSchedule — Deep Dive

What is PreferNoSchedule?

PreferNoSchedule is a **soft taint** — it tells Kubernetes:

“Try not to schedule pods here **if possible**, but it’s okay if there’s no better alternative.”

It is **advisory**, not mandatory.

When to Use PreferNoSchedule?

Use PreferNoSchedule when you want to:

- **Discourage** scheduling pods onto certain nodes
- Apply **soft separation** instead of hard enforcement
- Create **preference rules** without blocking workloads

Common use cases:

- Signal that a node is lightly reserved for a type of workload (e.g., batch jobs)
- Prefer that test pods avoid production nodes, but allow them during high resource pressure

-
- Let Kubernetes make the final decision if needed, while guiding its preference

Syntax and Example

Apply the taint:

```
kubectl taint nodes node-a workload=batch:PreferNoSchedule
```

This means:

👉 "Avoid scheduling here unless no other suitable nodes are available."

Pod Manifest (with or without toleration):

- ♦ Pod **without** toleration:
 - Kubernetes **will try** to avoid the tainted node.
 - **If no other nodes are available, it can still schedule here.**

♦ Pod **with** toleration:

tolerations:

```
- key: "workload"  
  operator: "Equal"  
  value: "batch"
```

effect: "PreferNoSchedule"

- - This doesn't force scheduling either.
 - It only says, "I'm okay running on this node if needed."



Behind the Scenes: How It Works

- The Kubernetes scheduler filters nodes with **NoSchedule** taints **out entirely** if not tolerated.
- But for **PreferNoSchedule**:
 - All nodes are kept in the candidate pool.
 - Nodes with **PreferNoSchedule** are given a **lower score**.
 - Scheduler still considers them but tries to choose others first.

This makes it ideal for **low-priority guidance**.

⚠️ Important Notes

- If all nodes are tainted with **PreferNoSchedule**, the pod **will be scheduled anyway**.
- It **does not protect** the node as strictly as **NoSchedule**.
- It is not suitable when **guaranteed isolation** is required.

✍️ Real-World Scenario

Problem:

You want to **prioritize critical apps** on certain nodes, but allow non-critical apps to run there **only if no other nodes are available**.

Solution:

Taint the preferred node group:

```
kubectl taint nodes prod-nodes  
critical=true:PreferNoSchedule
```

1. Deploy normal workloads without any toleration.

Result:

- Scheduler avoids **prod-nodes** for normal apps.

-
- But during heavy load, it will still schedule them if needed.

 **Advanced Tip: Combine with Pod Priorities**

`PreferNoSchedule` becomes powerful when used alongside **Pod PriorityClasses**.

- Critical pods → tolerate the taint and have higher priority
- Low-priority pods → no toleration, and get scheduled elsewhere

This allows you to implement **best-effort node preferences** based on workload importance.

🔥 NoExecute — Deep Dive

📌 What is NoExecute?

NoExecute is the **strictest and most aggressive taint effect** in Kubernetes.

It does two things:

1. **Prevents scheduling** of new pods on the tainted node (like NoSchedule)
2. **Evicts already running pods** from the node if they don't tolerate the taint

In short: NoExecute both **blocks new pods** and **kicks out existing ones** unless they have a toleration.

⚠️ Why It's Powerful (and Dangerous)

This effect is used when:

- You want to **forcefully evacuate** a node (e.g. for maintenance or failures)
- You want to **ensure only specific pods can ever run or stay** on that node
- You want to **respond dynamically to node health conditions** (e.g. taints from the node controller)

It's also used automatically by Kubernetes under certain system conditions — more on that below.



Behavior Summary

Pod State	Has Toleration	Effect
Not yet scheduled	✗	Pod won't be scheduled
Not yet scheduled	✓	Pod may be scheduled
Already running	✗	Pod gets evicted immediately
Already running	✓	Pod stays , until toleration duration expires (if set)



Syntax and Example

Taint a node:

```
kubectl taint nodes node-a role=critical:NoExecute
```

This taint means:

“Only pods that explicitly tolerate this taint will be allowed to run or stay.”

Pod With Infinite Toleration:**tolerations:**

```
- key: "role"  
  operator: "Equal"  
  value: "critical"  
  effect: "NoExecute"
```

This pod:

- Will be allowed to **stay running**
- Will be allowed to **be scheduled**

Pod With Temporary Toleration:**tolerations:**

```
- key: "role"  
  operator: "Equal"  
  value: "critical"  
  effect: "NoExecute"  
  
  tolerationSeconds: 60
```

This pod:

- Will be **evicted after 60 seconds**
- Useful for **graceful shutdown**, not instant eviction

Pod With No Tolerance:

- Will be **evicted immediately** if already running
- Will **not be scheduled** onto the tainted node at all



Kubernetes Itself Uses **NoExecute**

Kubernetes uses **NoExecute** taints internally, such as:

Node NotReady Taint

`node.kubernetes.io/not-ready:NoExecute`

1.

- Applied when a node becomes unreachable or unresponsive.
- Non-tolerating pods are evicted.

Node Unreachable Taint**node.kubernetes.io/unreachable:NoExecute**

2.

- Used when the kubelet stops reporting.
- Scheduler evicts pods if not tolerated.

In both cases, if you want your pods to **remain** during transient failures, add tolerations like:

tolerations:

```
- key: "node.kubernetes.io/unreachable"  
  operator: "Exists"  
  effect: "NoExecute"  
  
  tolerationSeconds: 300
```

This pod will remain for 5 minutes, allowing time for the node to recover.

 **Real-World Use Case****Scenario:**

You're managing a **dedicated compute node** in a data-intensive cluster.

Only long-running analytics workloads should run there, and **they must be drained quickly** when needed.

Solution:**Taint the node with NoExecute:**

```
kubectl taint nodes analytics-node batch=only:NoExecute
```

Analytics workloads tolerate it:**tolerations:**

```
- key: "batch"  
  value: "only"  
  effect: "NoExecute"
```

All other pods get instantly evicted if mistakenly scheduled.

🎯 When & Why Do We Use Taints and Tolerations in Production Kubernetes Clusters?

Taints and tolerations are **production-grade scheduling controls**. They are used to:

- Protect infrastructure components
- Isolate sensitive workloads
- Optimize resource usage
- Enforce multi-tenancy and security boundaries

Let's explore **real production scenarios** where these are used — with practical justification.



1. Protecting Critical Infrastructure Nodes

✳️ Situation:

You have dedicated nodes for:

- Control plane components

-
- Ingress controllers
 - Core DNS
 - Monitoring (Prometheus, Loki)

Why Taints?

You **don't want** regular app workloads or dev/test pods running on these nodes — they can starve system resources or impact reliability.

Example:

```
kubectl taint nodes node-monitor  
dedicated=monitoring:NoSchedule
```

Now only Prometheus pods that tolerate this taint will be scheduled here.

2. Node-Level Isolation for Different Environments

Situation:

You run **multiple environments** (dev, staging, prod) in the **same cluster**, but you want each to use its own node pool.

🎯 Why Taints?

- Prevent noisy neighbor issues
- Enforce performance separation
- Reduce risk of cross-env impact

✓ Example:

```
kubectl taint nodes prod-nodes env=prod:NoSchedule
```

Then prod-only workloads get a toleration like:

tolerations:

```
- key: "env"  
  value: "prod"  
  effect: "NoSchedule"
```

This ensures only prod workloads run on prod nodes.

🎮 3. Prioritizing GPU / Specialized Hardware Nodes

✖ Situation:

You have a pool of GPU-enabled nodes, or nodes with NVMe or FPGAs.

You want **only ML/AI workloads** to land on them — not regular backend pods.

⌚ Why Taints?

Prevents accidental usage of expensive nodes. Ensures hardware is available for critical jobs.

✓ Example:

```
kubectl taint nodes gpu-node gpu=true:NoSchedule
```

Then your ML training job can tolerate this taint:

tolerations:

- key: "gpu"
value: "true"
effect: "NoSchedule"



4. Evicting Pods from Unhealthy Nodes

Situation:

A node becomes **NotReady** or unreachable. You want pods to **automatically be evicted** and rescheduled elsewhere.

Why Taints?

Kubernetes automatically applies **NoExecute** taints:

- `node.kubernetes.io/unreachable`
- `node.kubernetes.io/not-ready`

You can add tolerations with **tolerationSeconds** to **delay eviction** and allow the node to recover:

tolerations:

```
- key: "node.kubernetes.io/not-ready"
  operator: "Exists"
  effect: "NoExecute"
  tolerationSeconds: 300
```

This keeps pods on the node for 5 minutes during transient issues.

5. Enforcing Pod Placement for Licensing / Compliance

Situation:

Certain applications (e.g., Oracle DB, SAP) must run on **licensed or pinned hardware nodes**.

Why Taints?

You can restrict those nodes and **only allow licensed apps** with matching tolerations.

Example:

```
kubectl taint nodes oracle-host license=oracle:NoSchedule
```

Only Oracle pods tolerate this taint.

6. Multi-Tenant Clusters — Tenant Isolation

Situation:

You run a **multi-team** or **multi-tenant** cluster.

You need to enforce soft or hard isolation between teams for:

- Compliance
- Cost visibility

-
- Debuggability

⌚ Why Taints?

- Reserve node groups for specific teams
- Prevent cross-team interference
- Combine with RBAC, network policies, and quotas

✓ Example:

```
kubectl taint nodes team-a-node tenant=team-a:NoSchedule
```

Team A workloads get a toleration:

tolerations:

```
- key: "tenant"  
  value: "team-a"  
  effect: "NoSchedule"
```



7. Graceful Drain or Rolling Update Strategy

Situation:

You need to **prepare a node for maintenance** (e.g., OS patching, autoscaling, draining).

Why Taints?

You can dynamically taint the node with **NoExecute** to begin gracefully evicting workloads.

Example:

```
kubectl taint nodes node-to-update  
upgrade=maintenance:NoExecute
```

Pods without tolerations are evicted.

Pods with **tolerationSeconds** get time to shut down cleanly.

Common Patterns for Using Taints & Toleration in Real-World Kubernetes Setups

These patterns are drawn from **actual production practices** across organizations that run **multi-tenant, resource-intensive, or security-sensitive** workloads.



Multi-Tenant Clusters (Team/Dept/Env Isolation)



Isolate workloads for different:

- Teams
- Environments (prod/stage/dev)
- Customers (in managed platforms)



Pattern:

Taint nodes per tenant/environment:

```
kubectl taint nodes node-1 tenant=team-a:NoSchedule
```

```
kubectl taint nodes node-2 tenant=team-b:NoSchedule
```

Workloads tolerate their taints:

tolerations:

```
- key: "tenant"  
  value: "team-a"  
  effect: "NoSchedule"
```

- Combine with:

- Namespaces per tenant
- ResourceQuotas
- NetworkPolicies
- RBAC

Benefit:

- Prevents tenant pods from running on wrong nodes
- Supports clear boundaries in a shared cluster



GPU / Specialized Hardware Workload Targeting

🎯 Goal:

Ensure only **ML/AI, data processing, or accelerated** workloads land on expensive hardware nodes.

✓ Pattern:

Taint GPU nodes:

```
kubectl taint nodes gpu-node accelerator=nvidia:NoSchedule
```

ML workloads tolerate:

tolerations:

```
- key: "accelerator"  
  value: "nvidia"  
  effect: "NoSchedule"
```

Combine with **NodeAffinity** to target node labels:

affinity:

```
nodeAffinity:  
  requiredDuringSchedulingIgnoredDuringExecution:  
    nodeSelectorTerms:
```

```
- matchExpressions:  
  - key: "kubernetes.io/instance-type"  
    operator: In  
    values:  
      - "gpu-optimized"
```

Benefit:

- Prevents non-GPU pods from consuming costly GPU nodes

System Component Isolation (Node Isolation)

Goal:

Prevent application pods from sharing nodes with system or critical services (e.g., CoreDNS, Ingress, Prometheus)

Pattern:

Taint **infra/system nodes**:

```
kubectl taint nodes system-node dedicated=system:NoSchedule
```

Allow only infra pods to tolerate it:

`tolerations:`

- `key: "dedicated"`
- `value: "system"`
- `effect: "NoSchedule"`
- Optionally combine with `nodeSelector` or `affinity`

Benefit:

- Guarantees that system services are isolated and not impacted by user workload spikes



Batch vs Interactive Workload Segregation

Goal:

Keep **non-critical**, **batch**, or **background jobs** from impacting **interactive** or **latency-sensitive** apps.

Pattern:

Taint low-priority batch nodes:

```
kubectl taint nodes batch-node  
workload=batch:PreferNoSchedule
```

Let batch workloads tolerate:

tolerations:

- key: "workload"

- value: "batch"

- effect: "PreferNoSchedule"

- Interactive workloads are **not scheduled** here by default

Benefit:

- Soft isolation with fallback flexibility
- Supports cost-effective autoscaling of batch workers



Controlled Node Maintenance / Eviction

Goal:

Drain nodes **gracefully** without disrupting services too quickly

Pattern:

During maintenance:

```
kubectl taint node node-x maintenance=planned:NoExecute
```

Let pods decide how long to tolerate:

tolerations:

- key: "maintenance"

- value: "planned"

- effect: "NoExecute"

- tolerationSeconds: 180

- Use **tolerationSeconds** to allow graceful shutdown



Benefit:

- Clean rolling updates or node patching
- Avoids chaotic evictions and service drops



Node Feature-Aware Scheduling



Goal:

Ensure that workloads land on nodes with specific **capabilities**, such as:

- AVX instruction sets
- SSD/NVMe
- High RAM

-
- Special kernel modules

 **Pattern:**

Taint high-capacity nodes:

```
kubectl taint nodes perf-node ssd=true:NoSchedule
```

App pods tolerate + use affinity:

tolerations:

```
- key: "ssd"
```

```
  value: "true"
```

```
  effect: "NoSchedule"
```

affinity:

nodeAffinity:

```
  requiredDuringSchedulingIgnoredDuringExecution:
```

nodeSelectorTerms:

```
      - matchExpressions:
```

```
        - key: "disk-type"
```

```
          operator: In
```

values:

```
            - "nvme"
```

 **Benefit:**

- Fine-grained workload placement
- Performance tuning via hardware-awareness



Taints and Tolerations in YAML — Structure & Syntax



💡 Taints — Defined on Nodes

Taints are added to **nodes** to **repel** pods unless they tolerate the taint.

✓ Add a taint to a node (CLI):

```
kubectl taint nodes <node-name> <key>=<value>:<effect>
```



🔍 Example:

```
kubectl taint nodes node1 env=prod:NoSchedule
```

This applies the following taint:

```
taints:
```

```
- key: "env"
```

```
  value: "prod"
```

```
  effect: "NoSchedule"
```



This means: “Only pods that tolerate this taint can be scheduled on this node.”



YAML Representation of Node With Taint

You can also define taints in a node YAML (for static files or custom tools):

```
apiVersion: v1
kind: Node
metadata:
  name: node1
spec:
  taints:
  - key: "env"
    value: "prod"
    effect: "NoSchedule"
```



Tolerations — Defined on Pods

Tolerations are defined **inside pod specs**, allowing pods to be scheduled onto nodes with matching taints.



Minimal YAML Structure

```
tolerations:
- key: "<taint-key>"
  operator: "Equal" | "Exists"
```

```
value: "<taint-value>"    # Optional if operator=Exists  
effect: "NoSchedule" | "PreferNoSchedule" | "NoExecute"  
tolerationSeconds: <int> # Only used with NoExecute
```

Example: Tolerate a **NoSchedule** Taint

tolerations:

```
- key: "env"  
operator: "Equal"  
value: "prod"  
effect: "NoSchedule"
```

 This pod can now be scheduled on nodes tainted with `env=prod:NoSchedule`.

Example: Tolerate Any **NoExecute** Taint for 60 Seconds

tolerations:

```
- key: "node.kubernetes.io/unreachable"  
operator: "Exists"  
effect: "NoExecute"
```

```
tolerationSeconds: 60
```

- ✓ This pod will be evicted after 60 seconds of the node becoming unreachable.

Example: Match Any Taint with **Exists**

```
tolerations:
```

```
- key: "env"  
  operator: "Exists"  
  effect: "NoSchedule"
```

- ✓ This tolerates any taint with key **env** and effect **NoSchedule**, regardless of value.

Complete Pod Spec Example

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: my-app
```

```
spec:
```

```
  containers:
```

```
    - name: app
```

```
      image: nginx
```

```
  tolerations:
```

```
    - key: "env"
```

```
      operator: "Equal"
```

```
      value: "prod"
```

```
      effect: "NoSchedule"
```

```
    - key: "maintenance"
```

```
      operator: "Equal"
```

```
      value: "planned"
```

```
      effect: "NoExecute"
```

```
  tolerationSeconds: 300
```

 This pod can:

- Be scheduled on `env=prod:NoSchedule` nodes

-
- Stay on `maintenance=planned :NoExecute` nodes for 5 minutes before eviction

Remove a Taint from a Node (CLI)

```
kubectl taint nodes <node-name> <key>:<effect>-
```

Example:

```
kubectl taint nodes node1 env=prod:NoSchedule-
```

 Removes the `env=prod :NoSchedule` taint

Real-Time Kubernetes Task:

Isolate Logging Workloads to Dedicated Nodes Using Taints & Toleration

Objective:

In a production-grade Kubernetes cluster (like your Kops-managed one), we want to **isolate logging agents** (e.g., Fluent Bit) to run **only** on specific nodes. These nodes should not run any other workloads.

Use Case:

- You have a DaemonSet (Fluent Bit) that collects logs from each node.
- You want only **dedicated logging nodes** to run the Fluent Bit pods.
- You'll taint these nodes so **no other workloads** get scheduled on them.
- Fluent Bit pods will tolerate these taints.

Steps to Implement

1. Label a Node as Logging Node

Pick one or more nodes you want to use only for logging.

```
kubectl label node ip-192-168-1-100 node-role=log
```

2. Taint the Node

```
kubectl taint nodes ip-192-168-1-100 logging=only:NoSchedule
```

 This tells Kubernetes:

"Only schedule pods that tolerate `logging=only:NoSchedule` on this node."

3. Create Fluent Bit DaemonSet with Toleration + Node Selector

Here's a **fully functional DaemonSet YAML** that will:

- Tolerate the taint
- Target only the labeled `log` node

-
- Deploy one pod per logging node

```
apiVersion: apps/v1

kind: DaemonSet

metadata:

  name: fluent-bit

  namespace: kube-logging

  labels:

    app: fluent-bit

spec:

  selector:

    matchLabels:

      app: fluent-bit

  template:

    metadata:

      labels:

        app: fluent-bit

  spec:

    tolerations:

      - key: "logging"
```

```
operator: "Equal"
value: "only"
effect: "NoSchedule"

nodeSelector:
  node-role: log

containers:
  - name: fluent-bit
    image: fluent/fluent-bit:2.2

resources:
  limits:
    memory: "200Mi"
    cpu: "100m"
  requests:
    memory: "100Mi"
    cpu: "50m"

volumeMounts:
  - name: varlog
```

```
        mountPath: /var/log

        - name: varlibdockercontainers
          mountPath: /var/lib/docker/containers
          readOnly: true

  volumes:
    - name: varlog
      hostPath:
        path: /var/log

    - name: varlibdockercontainers
      hostPath:
        path: /var/lib/docker/containers
```

4. Apply the DaemonSet

```
kubectl create namespace kube-logging

kubectl apply -f fluent-bit-daemonset.yaml
```

✓ 5. Validate Everything

🔍 Check the taint:

```
kubectl describe node ip-192-168-1-100 | grep Taints
```

Should output:

```
Taints: logging=only:NoSchedule
```

🔍 Check that Fluent Bit runs only on the logging node:

```
kubectl get pods -n kube-logging -o wide
```

You should see pods running **only on ip-192-168-1-100**.

🔍 Check that no other workloads are scheduled there:

```
kubectl get pods --all-namespaces -o wide | grep  
ip-192-168-1-100
```

Should show **only Fluent Bit**, nothing else.

Why This Matters in Production

- Keeps logging isolated from noisy neighbor workloads
- Ensures log agents are not starved for resources
- Enables fine-grained control over resource sizing and monitoring
- Used in large-scale infra at AWS, GCP, and enterprise setups

Bonus Tip: Automate with Kops

If you're using Kops, you can define taints in the instance group like this:

`taints:`

```
- key: "logging"  
  value: "only"  
  effect: "NoSchedule"
```

This will taint the nodes automatically during cluster creation or rolling update.

Real-Time Kubernetes Task #2:

Graceful Node Draining for Maintenance Using **NoExecute** Taints

Objective:

When a node is going into **maintenance (patching, reboot, draining)**, you want to **gracefully evict workloads** — but **only those that don't tolerate the taint**.

This allows:

- Critical pods to stay temporarily
- Others to evict cleanly
- Controlled disruption in production

Use Case:

- You need to **restart or upgrade** a node without killing all pods instantly.
- You apply a **temporary taint** to evict non-critical workloads.
- Some critical workloads have a **tolerationSeconds** buffer to remain briefly and then gracefully shut down.



Implementation Steps



1. Choose a Node for Maintenance

Example:

```
kubectl get nodes
```

Pick a node (say, **ip-192-168-1-80**) that needs to be restarted or upgraded.



2. Apply a Temporary Maintenance Taint (NoExecute)

```
kubectl taint nodes ip-192-168-1-80  
maintenance=planned:NoExecute
```

This will **evict all pods** that **don't tolerate** this taint.



3. Deploy a Pod That Can Stay Temporarily (With tolerationSeconds)

Here's an example of a **critical pod** that can tolerate the maintenance taint for **300 seconds (5 minutes)**:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
name: payment-processor
```

```
labels:
```

```
  app: payment
```

```
spec:
```

```
  containers:
```

```
    - name: main
```

```
      image: nginx
```

```
  tolerations:
```

```
    - key: "maintenance"
```

```
      operator: "Equal"
```

```
      value: "planned"
```

```
      effect: "NoExecute"
```

```
  tolerationSeconds: 300
```

 This pod:

- Will **not be evicted immediately**
- Will be given **5 minutes** to finish gracefully (e.g., draining a queue, flushing a cache)

✓ 4. Deploy a Pod That Gets Evicted Immediately

```
apiVersion: v1
kind: Pod
metadata:
  name: web-frontend
  labels:
    app: frontend
spec:
  containers:
    - name: web
      image: nginx
```

🚫 This pod has **no tolerations**, so when the node gets tainted, it is **evicted immediately**.

✓ 5. Observe Eviction Behavior

Use:

```
kubectl get pods -o wide
kubectl describe pod payment-processor
```

✓ You'll see:

- **web-frontend** is terminated immediately
- **payment-processor** runs for 5 minutes, then is evicted

✓ 6. Cleanup Taint After Maintenance

Once the node is patched, remove the taint:

```
kubectl taint nodes ip-192-168-1-80  
maintenance=planned:NoExecute-
```

🧠 Why This Is Production-Ready

- Used for **graceful upgrades and maintenance windows**
- Ensures **high availability** for critical services
- Prevents **violent pod termination**
- Gives control to platform engineers while respecting app SLOs

 **Bonus Tip: Use with Cluster Autoscaler**

Cluster autoscaler respects **NoExecute** taints and **doesn't scale down** nodes that have tainted but non-evictable pods.

Real-Time Kubernetes Task #3:

Isolating GPU Workloads on Specialized Nodes Using Taints and Tolerations

Objective:

You want to schedule **machine learning (ML) or GPU-based workloads** (like TensorFlow, PyTorch, or inference APIs) **only on GPU-enabled nodes** in your cluster.

Other workloads should **never run on GPU nodes** to preserve resources and reduce scheduling conflicts.

Use Case:

- You provision dedicated GPU nodes in your Kops cluster (using a separate instance group).
- You taint those nodes to prevent regular pods from landing there.
- You use tolerations on GPU-related pods so **only those can be scheduled**.



Implementation Steps



1. Label and Taint the GPU Nodes

Let's assume a node with GPU is: **ip-192-168-1-90**

```
kubectl label node ip-192-168-1-90 node-type=gpu  
kubectl taint node ip-192-168-1-90  
accelerator=nvidia:NoSchedule
```

 Now, only pods with the right toleration can be scheduled on this GPU node.



2. Create a GPU-Based Workload (With Toleration)

Here's a **sample ML inference pod** that uses NVIDIA GPU and tolerates the taint.

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: image-classifier  
spec:  
  containers:  
    - name: classifier  
      image: nvcr.io/nvidia/tensorflow:23.06-tf2-py3
```

```
resources:
```

```
limits:
```

```
  nvidia.com/gpu: 1
```

```
tolerations:
```

```
- key: "accelerator"  
  operator: "Equal"  
  value: "nvidia"  
  effect: "NoSchedule"
```

```
nodeSelector:
```

```
  node-type: gpu
```

 Notes:

- `resources.limits.nvidia.com/gpu: 1` requests 1 GPU
- The toleration matches the taint on the node
- `nodeSelector` ensures it's scheduled only on GPU-labeled nodes

✓ 3. Deploy the Pod

```
kubectl apply -f image-classifier.yaml
```

Check:

```
kubectl get pod image-classifier -o wide
```

✓ Should be running on **ip-192-168-1-90** (the GPU node)

✓ 4. Try to Schedule a Regular Pod Without Toleration

```
apiVersion: v1
kind: Pod
metadata:
  name: web-frontend
spec:
  containers:
    - name: web
      image: nginx
```

```
kubectl apply -f web-frontend.yaml
```

This pod **will not get scheduled** on the GPU node because:

- It doesn't tolerate the `accelerator=nvidia:NoSchedule` taint
- The scheduler avoids GPU nodes

5. Validate the Setup

```
kubectl get pods -o wide
```

```
kubectl describe node ip-192-168-1-90 | grep Taint
```

You'll see:

```
Taints: accelerator=nvidia:NoSchedule
```

And only GPU pods are allowed on it.

Why This Pattern Is Important in Production

- Prevents **non-GPU workloads** from wasting GPU resources
- Enforces **workload affinity** in heterogeneous node pools
- Useful for **ML workloads, video processing, AR/VR, AI inference APIs**
- Supports efficient **resource isolation and scaling**



Kops Integration Tip:

In your Kops instance group for GPU nodes, add this:

taints:

```
- key: "accelerator"  
  value: "nvidia"  
  effect: "NoSchedule"
```

labels:

```
node-type: gpu
```

This automatically taints and labels nodes during cluster update.

Real-Time Kubernetes Task #4:

Isolating Team Workloads in a Multi-Tenant Cluster Using Taints and Toleration

Objective:

In a **multi-team (multi-tenant)** Kubernetes cluster managed via Kops, ensure that workloads from different teams **run on dedicated node pools**. Prevent "noisy neighbor" issues, resource competition, or accidental deployments across teams.

Use Case:

- Teams like **team-a**, **team-b**, and **platform** share the same cluster.
- Each team has its own **dedicated node group** via Kops.
- Taints are used to **repel unrelated workloads** from those nodes.
- Tolerations ensure only the owning team's workloads get scheduled on the right nodes.



Implementation Steps

1. Set Up Node Isolation via Kops (Taint + Label)

Assume a dedicated instance group for `team-a`. In the instance group YAML (via `kops edit ig team-a`):

`taints:`

```
- key: "dedicated"  
  value: "team-a"  
  effect: "NoSchedule"
```

`labels:`

```
node-pool: team-a
```

After updating:

```
kops update cluster --yes
```

```
kops rolling-update cluster --yes
```

This:

- Taints all nodes in this group with `dedicated=team-a:NoSchedule`
- Labels them with `node-pool=team-a`

✓ Only pods that tolerate `dedicated=team-a` will be scheduled there.

✓ 2. Deploy Team-A Workloads With Toleration + NodeSelector

Here's a sample Deployment for `team-a` that is correctly configured:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: team-a-api
  namespace: team-a
spec:
  replicas: 2
  selector:
    matchLabels:
      app: team-a-api
  template:
    metadata:
      labels:
        app: team-a-api
  spec:
    tolerations:
```

```
- key: "dedicated"
  operator: "Equal"
  value: "team-a"
  effect: "NoSchedule"

nodeSelector:
  node-pool: team-a

containers:
  - name: api
    image: teama/api:latest

resources:
  requests:
    memory: "256Mi"
    cpu: "250m"
  limits:
    memory: "512Mi"
    cpu: "500m"
```

✓ 3. Try Deploying Without Toleration (Simulate Misconfiguration)

Let's say someone on **team-b** mistakenly tries to deploy to **team-a**'s node group:

```
apiVersion: v1
kind: Pod
metadata:
  name: rogue-pod
  namespace: team-b
spec:
  containers:
  - name: nginx
    image: nginx
  nodeSelector:
    node-pool: team-a
```

🔍 Result:

```
kubectl describe pod rogue-pod -n team-b
```

– You'll see:

node(s) had taint {dedicated=team-a: NoSchedule}, that the pod didn't tolerate

✓ This protects node isolation.

✓ 4. Validation

🔍 Check taints on the node:

```
kubectl describe node <node-name> | grep Taint
```

🔍 Confirm pod placement:

```
kubectl get pods -o wide -n team-a
```

Should show **team-a-api** pods only on **node-pool: team-a** nodes.

Why This Pattern Works in Production

- Enforces **team-level workload isolation**
- Prevents **accidental cross-scheduling**
- Works well in **shared clusters**
- Scales with your **infrastructure-as-code via Kops**

Used heavily by:

- Platform teams in mid-large orgs
- SaaS companies with shared Kubernetes clusters
- Cost optimization: isolate by environment (prod/dev/staging)



Bonus Tip: Namespace + RBAC + Tainting = Real Isolation

- Namespace per team (✓)
- RBAC scoped to namespace (✓)
- Node taints per team (✓)
- ResourceQuota and LimitRange (🔒)
- NetworkPolicy (optional 🔒)

Put together = multi-tenant SaaS-grade architecture.

Real-Time Kubernetes Task #5:

Blue-Green Deployment Separation Using Taints and Toleration

Objective:

Implement **blue-green deployments** in Kubernetes using **taints and tolerations** to isolate traffic between two environments during rollout, testing, and production cutover.

This method allows:

- **Two versions of the app to run simultaneously**
- Controlled traffic switching
- Full environment isolation before cutover
- Seamless rollback

Use Case:

- You run both **blue** (live) and **green** (new) environments in parallel.
- Blue and Green workloads **run on separate node pools**.

-
- During testing, only internal teams access green.
 - After validation, you switch traffic to green (e.g., via Ingress, DNS, or service label changes).

Implementation Steps

1. Provision Dedicated Node Groups (Kops)

Create 2 instance groups via Kops:

- `blue-group` → Live traffic
- `green-group` → Staging/new version

In `blue-group`, add:

`labels:`

`deployment-color: blue`

`taints:`

`- key: "env"`

`value: "blue"`

`effect: "NoSchedule"`

In green-group, add:

labels:

```
deployment-color: green
```

taints:

```
- key: "env"
```

```
value: "green"
```

```
effect: "NoSchedule"
```

Apply and roll the cluster:

```
kops update cluster --yes
```

```
kops rolling-update cluster --yes
```

2. Deploy the Blue Version of the App

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
name: myapp-blue
```

```
spec:
```

```
replicas: 2
```

```
selector:  
  
  matchLabels:  
  
    app: myapp  
  
    color: blue  
  
template:  
  
  metadata:  
  
    labels:  
  
      app: myapp  
  
      color: blue  
  
spec:  
  
  nodeSelector:  
  
    deployment-color: blue  
  
  tolerations:  
  
    - key: "env"  
  
      operator: "Equal"  
  
      value: "blue"  
  
      effect: "NoSchedule"  
  
  containers:  
  
    - name: app
```

```
image: myorg/myapp:1.0.0
```

```
ports:
```

```
- containerPort: 80
```

Pods only run on **blue** tainted nodes.

3. Deploy the Green Version (New Release Candidate)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: myapp-green
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      app: myapp
```

```
      color: green
```

```
  template:
```

```
    metadata:
```

```
labels:  
  app: myapp  
  color: green  
  
spec:  
  nodeSelector:  
    deployment-color: green  
  
  tolerations:  
    - key: "env"  
      operator: "Equal"  
      value: "green"  
      effect: "NoSchedule"  
  
  containers:  
    - name: app  
      image: myorg/myapp:2.0.0  
  
  ports:  
    - containerPort: 80
```

 Pods for green version run **exclusively** on green nodes.

✓ 4. Route Traffic Strategically

Use a Service or Ingress to control traffic:

Initial Setup (Blue is Live):

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: myapp
    color: blue
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Switch to Green (Post Validation):

Update the selector:

`spec:`

`selector:`

`app: myapp`

`color: green`

Or if using Ingress (e.g., with nginx or alb), point the path/backend to green version.

5. Optional Rollback Strategy

If green version fails tests or SLOs:

- Simply revert the service or Ingress back to blue
- Delete or scale down green Deployment

No need to re-schedule or rollback nodes — isolation is already in place.

Validate

Check pods and nodes:

```
kubectl get pods -o wide
```

```
kubectl describe node <node-name>
```

You'll observe:

- Pods only scheduled to the intended tainted nodes
- No workload interference between environments



Why This Pattern Works in Production

- True **environment-level isolation**
- **Quick rollback** by switching traffic
- Zero-downtime deployment pattern
- Excellent for regulated systems, canary releases, A/B testing, or client-specific upgrades



Troubleshooting: Taints and Tolerations in Kubernetes

✖ Problem 1: Pod is not getting scheduled (Pending state)

🔍 Symptom:

```
kubectl get pod <pod-name>
```

STATUS: Pending

🔍 Diagnostic Command:

```
kubectl describe pod <pod-name>
```

✓ Typical Output:

Events:

```
Warning FailedScheduling ... 0/3 nodes are available: 3
node(s) had taint {dedicated=team-a: NoSchedule}, that the
pod didn't tolerate
```

 **Root Cause:**

Pod is trying to land on tainted nodes but doesn't **have matching tolerations**.

 **Solution:**

Add the correct toleration to the Pod spec:

tolerations:

```
- key: "dedicated"  
  operator: "Equal"  
  value: "team-a"  
  
effect: "NoSchedule"
```

 **Problem 2: Pod has toleration, but still not scheduled** **Symptom:**

Pod has a toleration, but it's still pending and not running on expected node.

 **Check NodeSelector or Affinity:**

```
kubectl describe pod <pod-name> | grep -A 10 Tolerations
```

Also check:

```
kubectl describe node <node-name> | grep Labels
```

 **Root Cause:**

- Toleration is present 
- But Pod has **nodeSelector** or **affinity** that doesn't match the tainted node's labels 

 **Solution:**

Make sure the **nodeSelector** matches the node's label:

nodeSelector:

```
node-pool: team-a
```

Or check **affinity/preferredDuringScheduling** rules if using them.

 **Problem 3: Pod keeps restarting or getting evicted** **Symptom:**

Pod shows **Evicted**, or restarts frequently.

 **Check Events or Describe Pod:**

```
kubectl get pod <pod-name>
```

```
kubectl describe pod <pod-name>
```

 **Root Cause:**

Node has a **NoExecute** taint — and the pod didn't tolerate it.

So it's being evicted as soon as it lands.

 **Solution:**

Add **NoExecute** toleration with optional **tolerationSeconds**:

tolerations:

```
- key: "critical"  
  operator: "Equal"  
  value: "true"  
  effect: "NoExecute"  
tolerationSeconds: 60
```

This keeps the pod for 60 seconds before eviction.

 **Problem 4: Taint removed, but pods still not scheduled** **Symptom:**

Even after removing taint, pods don't get scheduled to the node.

 **Double Check Taint Removal:**

```
kubectl describe node <node-name>
```

 **Root Cause:**

- Taint might have been removed 
- But other **scheduling conditions** (like disk pressure, node unschedulable, etc.) are blocking pod placement

 **Solution:**

Check full node status:

```
kubectl get nodes
```

```
kubectl describe node <node-name>
```

Look for:

- **SchedulingDisabled**
- DiskPressure
- PIDPressure
- MemoryPressure

Problem 5: Pod lands on wrong node despite taints

Symptom:

You applied a taint on the node, but a pod still lands on it.

Check Taint Effect:

```
kubectl describe node <node-name>
```

Root Cause:

- Taint might be using **PreferNoSchedule**, which is a **soft** taint (scheduler tries to avoid but not guaranteed).
- Or pod has a **wildcard toleration**, which tolerates all taints.

Solution:

 Use **NoSchedule** instead of **PreferNoSchedule** for hard enforcement.

Also, avoid wildcard tolerations unless necessary:

tolerations:

```
- operator: "Exists"
```

This tolerates all taints — use carefully!

✖ Problem 6: YAML syntax issues when defining taints/tolerations

⚠ Common Mistakes:

- Wrong indentation
- **effect** missing
- Using **tolerate** instead of **tolerations**
- Misplaced under metadata

✓ Correct Toleration Block:

tolerations:

```
- key: "example"  
  operator: "Equal"  
  value: "test"  
  
  effect: "NoSchedule"
```

Use a linter or validate your YAML:

```
kubectl apply --dry-run=client -f myfile.yaml
```



Commands Summary for Debugging

Task	Command
List taints on node	<code>kubectl describe node <node></code>
Describe pod for scheduling info	<code>kubectl describe pod <pod></code>
View unschedulable pods	<code>kubectl get pods --field-selector status.phase=Pending</code>
Dry-run to validate YAML	<code>kubectl apply --dry-run=client -f file.yaml</code>



Best Practices to Avoid Taint Issues

- Always pair **taints** with **labels** and use **nodeSelector** or **affinity** to schedule accurately.
- Avoid overusing **NoExecute** unless you need eviction behavior.
- Set alerts for **Pending pods** via monitoring tools (Prometheus, Datadog, etc.).

-
- Document every taint in your infra-as-code (e.g., Kops instance groups or Terraform).



Best Practices for Taints and Tolerations in Kubernetes

⌚ 1. Use Taints to Enforce Node Roles — Not Replace Scheduling Policies

- Taints **should not be used alone** for workload placement.
- Always combine with:
 - `nodeSelector`
 - `nodeAffinity`
 - `node labels`



Example:

`nodeSelector:`

`workload-type: gpu`

`tolerations:`

`- key: "gpu"`

`operator: "Equal"`

```
value: "true"  
effect: "NoSchedule"
```

 This gives **hard isolation** and **explicit scheduling**, not accidental pod placement.

2. Avoid Wildcard Tolerations Unless Absolutely Required

tolerations:

- operator: "Exists"

- This tolerates **any taint**, which could lead to **pods landing on restricted nodes** (like control-plane or spot-only pools).
- Use only in **infrastructure-level DaemonSets** like `fluentd`, `cilium`, `kube-proxy`, etc.

 Instead, use **scoped tolerations** with key/value pairs.

⌚ 3. Label Every Tainted Node Group and Use NodeSelector

- Taints block scheduling 
- Tolerations allow it 
- But without a label-selector, the scheduler may try to place pods across unrelated tainted pools.

 Pair taints with **labels** and schedule pods precisely:

labels:

```
node-group: platform-nodes
```

nodeSelector:

```
node-group: platform-nodes
```

⌚ 4. Use **NoExecute** With Care

- **NoExecute** evicts pods **immediately or after tolerationSeconds**.
- Use in **health-related scenarios** like:
 - Node fails probe

- Node enters disk-pressure state

 Add `tolerationSeconds` to avoid instant eviction during minor blips:

`tolerations:`

- `key: "node.kubernetes.io/not-ready"`

`effect: "NoExecute"`

`tolerationSeconds: 60`

5. Isolate Critical Workloads With Taints

For example:

- Platform-level services
- Logging agents
- System-level DaemonSets
- High-priority APIs

 Taint their dedicated nodes and use strict tolerations.

`taints:`

- `key: "dedicated"`

`value: "platform"`

```
effect: "NoSchedule"
```

⌚ 6. Default Node Pools Should Remain Untainted for General Workloads

- Keep your “default” or “generic” node pool **untainted**.
- Only apply taints where **isolation is required** (e.g., GPU, staging, infra).

This avoids broken scheduling for non-tolerating pods by default.

⌚ 7. Document and Standardize Taints in GitOps/IaC

- Taints should live in:
 - **kops** instance group YAMLs
 - **terraform** configs
 - **Cluster API** specs

✓ Track taints the same way you do labels, annotations, node selectors, etc.

taints:

```
- key: "workload"
```

```
value: "ml"
```

```
effect: "NoSchedule"
```

⌚ 8. Validate Taints and Tolerations During CI

Use tools like:

- [kubeval](#) or [kubeconform](#)
- [OPA Gatekeeper](#) policies
- Admission controllers

✓ Example Gatekeeper policy: prevent wildcard tolerations in dev workloads.

⌚ 9. Use Monitoring for Pending Pods and Evictions

Track:

- [Pending](#) pod counts
- [Evicted](#) pod reasons
- [taint_toleration_mismatch](#) in logs or metrics

 Alert if:

- Pods are pending for > X seconds
- Node taints unexpectedly increase (e.g., due to autoscaler)

10. Use Taints for Spot Node Isolation

For cost-effective, fault-tolerant workloads:

`taints:`

```
- key: "spot"  
  value: "true"  
  effect: "NoSchedule"
```

 Schedule only retryable or stateless workloads with this toleration.



Bonus Best Practice: Taint the Control Plane

In most Kubernetes setups (including Kops), control-plane nodes come pre-tainted:

```
node-role.kubernetes.io/control-plane:NoSchedule
```



Never remove this taint unless:

- You're doing single-node testing (NOT recommended in prod)
- You **intentionally** want to run workloads there (with toleration)



Node Labels and Selectors in Kubernetes



What are Node Labels?

Node labels are simple key-value pairs attached to Kubernetes nodes.

They are used to **classify** and **group** nodes based on attributes like:

- Instance type
- Environment (dev/prod)
- Zone or region
- GPU/CPU specialization
- Node group name
- Spot vs on-demand



Labels are **immutable per node lifecycle**, but can be changed dynamically.

```
kubectl label node ip-172-20-30-100 workload=gpu
```

 Syntax Example

labels:

```
node-role.kubernetes.io/worker: ""
instance-type: "c5.large"
environment: "production"
workload: "ml"
```

 How Do We Use Labels?

They are **used by pods to request specific nodes** via:

- **nodeSelector**
- **nodeAffinity**

This ensures **workload placement** aligns with node capability or purpose.

nodeSelector: Simple and Common

What is nodeSelector?

`nodeSelector` is the **simplest way** to bind a pod to a node with specific labels.

`spec:`

`nodeSelector:`

`workload: ml`

`environment: production`

This will schedule the pod **only** on nodes that have both:

- `workload=ml`
- `environment=production`

Real Example: Run a Pod on GPU Nodes

Let's assume nodes in your cluster are labeled like:

```
kubectl label node ip-172-20-50-10 gpu=true
```

Now create a pod with:

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-pod
spec:
  containers:
  - name: app
    image: tensorflow/tensorflow:latest-gpu
  nodeSelector:
    gpu: "true"
```

 This ensures the pod runs only on GPU-labeled nodes.

 **Use Case Scenarios**

Use Case	Label	Selector Example
Environment-specific	<code>env=prod / env=dev</code>	<code>nodeSelector: { env: prod }</code>
Node group isolation	<code>node-pool=infra</code>	<code>nodeSelector: { node-pool: infra }</code>
Architecture matching	<code>beta.kubernetes.io/arch</code>	<code>nodeSelector: { arch: amd64 }</code>
Region/zone targeting	<code>topology.kubernetes.io/zone</code>	<code>nodeSelector: { ... }</code>

 **Node Affinity (Advanced Alternative)**

If you want more expressive rules (like OR/AND logic), use **Node Affinity** instead of `nodeSelector`.

`affinity:`

`nodeAffinity:`

`requiredDuringSchedulingIgnoredDuringExecution:`

`nodeSelectorTerms:`

`- matchExpressions:`

```
- key: environment  
operator: In  
values:  
- production  
- staging
```

 More flexibility than `nodeSelector`, especially in multi-zone, multi-tenant environments.

Commands for Node Labels

View Node Labels

```
kubectl get nodes --show-labels
```

Add Label to a Node

```
kubectl label node <node-name> zone=us-east-1a
```

Remove Label from Node

```
kubectl label node <node-name> zone-
```

Best Practices for Node Labels & Selectors

Label Strategy

- Use **consistent naming** conventions (e.g., `env`, `node-pool`, `zone`, `workload`).
- Use **labels** for capabilities, and **taints/tolerations** for restrictions.
- Document every label in GitOps/IaC (like in Kops instance group YAMLS).

Don't Rely on Dynamic Labels for Scheduling

- Dynamic labels (e.g., added at runtime via script) might cause flaky scheduling.
- Label nodes at **provision time** (via `Kops`, Terraform, etc.).

✓ YAML Summary Example (Pod + Node)

```
# Assume this node has:  
  
# labels: workload=ml, env=prod  
  
apiVersion: v1  
  
kind: Pod  
  
metadata:  
  
    name: ml-job  
  
spec:  
  
    containers:  
  
        - name: app  
  
            image: ml-image:latest  
  
    nodeSelector:  
  
        workload: ml  
  
        env: prod  
  
    tolerations:  
  
        - key: "dedicated"  
  
            operator: "Equal"  
  
            value: "ml"
```

```
effect: "NoSchedule"
```

 This pod will only run on nodes that match the labels AND tolerate the taint — perfect for isolated workloads like ML, platform services, or GPU jobs.



Taints & Tolerations + Node Labels & Selectors — Complete Concept



Understanding the Roles

Feature	Purpose
Labels	Add metadata to nodes (e.g., type, zone, purpose)
Selectors	Used in pods to choose nodes with specific labels
Taints	Applied to nodes to repel pods unless they tolerate them
Tolerations	Applied to pods so they can be scheduled on tainted nodes



How They Work Together

- **Labels & nodeSelectors:** Used to **target** specific node(s).
- **Taints & tolerations:** Used to **restrict** which pods can run on those nodes.

-  Use **both** to avoid accidental scheduling or unwanted behavior.

Real-World Example Scenario: Dedicated ML Workloads on GPU Nodes

1. Label the GPU Nodes

```
kubectl label node ip-10-0-0-5 workload=gpu
```

Now the node has:

labels:

```
workload: gpu
```

2. Taint the GPU Node

```
kubectl taint nodes ip-10-0-0-5 dedicated=gpu:NoSchedule
```

Now the node has:

taints:

```
- key: "dedicated"
```

```
value: "gpu"
```

```
effect: "NoSchedule"
```

 This prevents any pod from being scheduled **unless it tolerates** this taint.

3. Deploy Pod That Matches Label + Tolerates Taint

```
apiVersion: v1
kind: Pod
metadata:
  name: ml-job
spec:
  containers:
    - name: trainer
      image: tensorflow/tensorflow:latest-gpu
  nodeSelector:
    workload: gpu
  tolerations:
    - key: "dedicated"
      operator: "Equal"
      value: "gpu"
```

```
effect: "NoSchedule"
```

What Happens Here?

1. Kubernetes **finds nodes** with `workload=gpu` — because of the `nodeSelector`.
2. Scheduler sees the **taint** on those nodes: `dedicated=gpu:NoSchedule`.
3. Pod has a **matching toleration** — so it is allowed to schedule on those nodes.
4. Pod runs only on GPU nodes and stays isolated from the rest of the cluster.

Why Use Both Together?

Using **only labels** means **any pod** can try to run on the node — not safe for production isolation.

Using **only taints** means **any pod with a wildcard toleration** may sneak in — not precise.

Best Practice:

- Use **labels + nodeSelector** to **target** node

-
- Use **taints + tolerations** to control access



Example: Misuse Without Selector

If the pod had only toleration and no selector:

tolerations:

```
- key: "dedicated"  
  operator: "Equal"  
  value: "gpu"  
  effect: "NoSchedule"
```

✗ Scheduler might still place this pod **on any node with that taint**, even if it's not intended for GPU workloads.

This is risky in **shared or multi-tenant clusters**.

Wrapping Up

Thank you so much for taking the time to read this documentation on **Kubernetes Taints, Tolerations, Node Labels & Selectors**. 

We've gone beyond just definitions — diving into how these features work together to **control, optimize, and isolate workload scheduling** in real-world production clusters. From foundational concepts to practical implementation using **kops**, you now have the tools to confidently apply these patterns in your own Kubernetes environments.



What's Next?

This doc is part of a larger series focused on hands-on, production-grade DevOps knowledge across:

-  **AWS Infrastructure**
-  **Ansible Automation**
-  **Kubernetes Operations**
-  **Real-world Project Scenarios**
-  **Scalable Cloud-Native Practices**

If you found this helpful, make sure to check out my other write-ups — and stay tuned for even more practical, well-structured docs coming very soon.

🙏 Thank You

Once again, thank you for being here and for supporting my DevOps content journey.

Whether you're just starting out or running production workloads, I hope this guide gave you something valuable.

📌 **Follow me on [LinkedIn] and [GitHub]** for more upcoming deep dives, tutorials, and end-to-end DevOps scenarios — all with a strong focus on practical knowledge and real-world relevance.

👋 Until next time — stay curious, keep experimenting, and happy deploying!

More docs are on the way. 🔥